

Cody Hubbard 004843389

CSM146, Winter 2018

Problem Set 2

Problem 1

(a) Solution: For AND:

$\theta(x_1, x_2) = \left(\begin{bmatrix} 1 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} - 1 \right)$ this can be tested by plugging in values for x_1 and x_2 .

$\theta(1, 1)$	1	+
$\theta(1, -1)$	-1	-
$\theta(-1, 1)$	-1	-
$\theta(-1, -1)$	-3	-

Another valid perception is $\theta_2(x_1, x_2) = \left(\begin{bmatrix} 1 & \frac{1}{2} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} - 1 \right)$ which can be shown to satisfy AND as

$\theta(1, 1)$	$= \frac{1}{2}$	+
$\theta(1, -1)$	$= -\frac{1}{2}$	-
$\theta(-1, 1)$	$= -\frac{3}{2}$	-
$\theta(-1, -1)$	$= -\frac{5}{2}$	-

well

(b) Solution: No two input perception exists to compute the XOR function. This is because of the way the sets of points that map to -1 and 1 intersect make it impossible to linearly separate them.

Problem 2

Solution: given $J(\theta) = -\sum_{n=1}^N [y_n \log(h_\theta(x_n)) + (1-y_n) \log(1-h_\theta(x_n))]$ and $h_\theta = \sigma(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$

Looking at that function you can see that $\frac{\partial J}{\partial \theta}$ comes down to $\log(h_\theta(x_n))'$ and $\log(1-h_\theta(x_n))'$

$$\log(h_\theta(x_n))' = \frac{1}{h_\theta(x_n)} \cdot h_\theta(x_n)'$$

where

$$\begin{aligned} h(x_n)' &= \left(\frac{1}{1 + e^{-\theta^T x}} \right)' \\ &= \frac{x e^{-\theta^T x}}{(1 + e^{-\theta^T x})^2} \end{aligned}$$

thus

$$\begin{aligned}
\log(h_\theta(x_n))' &= \frac{1}{h_\theta(x_n)} \cdot h_\theta(x_n)' \\
&= \left(1 + e^{-\theta^T x}\right) \cdot \frac{x e^{-\theta^T x}}{(1 + e^{-\theta^T x})^2} \\
&= \frac{x e^{-\theta^T x}}{1 + e^{-\theta^T x}} \\
&= x e^{-\theta^T x} \cdot h_\theta(x_n)
\end{aligned}$$

and

$$\log(1 - h_\theta(x_n))' = x e^{-\theta^T x}$$

and finally

$$\frac{\partial J}{\partial \theta} = - \sum_{n=1}^N [y_n (x e^{-\theta^T x} \cdot h_\theta(x_n)) + (1 - y_n)(x e^{-\theta^T x})]$$

Problem 3

(a) Solution: Given $J(\theta_0, \theta_1) = \sum_{n=1}^N w_n(\theta_0 + \theta_1 x_{n,1} - y_n)^2$

$$\begin{aligned}
\frac{\partial J}{\partial \theta_0} &= \sum_{n=1}^N 2 \cdot w_n(\theta_0 + \theta_1 x_{n,1} - y_n) \\
\frac{\partial J}{\partial \theta_1} &= \sum_{n=1}^N 2 \cdot x_n \cdot w_n(\theta_0 + \theta_1 x_{n,1} - y_n)
\end{aligned}$$

(b) Solution: Take the sum so that the full vectors for X, Y, W , and θ are formed, then

$$\begin{aligned}
J(\theta_0, \theta_1) &= (X\theta - Y)^T W(X\theta - Y) \\
&= (X\theta - Y)^T W(X\theta - Y) \\
&= (\theta^T X^T - Y^T) W(X\theta - Y) \\
&= (\theta^T X^T - Y^T)(WX\theta - WY) \\
&= (\theta^T X^T WX\theta - Y^T WX\theta - \theta^T X^T WY - Y^T WY)
\end{aligned}$$

and now take the derivative and set equal to zero

$$\begin{aligned}
\frac{\partial J}{\partial \theta} &= (X^T WX\theta - Y^T WX - X^T WY) \\
0 &= X^T WX\theta - Y^T WX - X^T WY \\
(X^T WX)\theta &= Y^T WX + X^T WY \\
\theta &= (X^T WX)^{-1} Y^T WX + (X^T WX)^{-1} X^T WY \\
\theta &= (X^T WX)^{-1} Y^T WX + X^{-1} Y
\end{aligned}$$

Problem 4

(a) Solution: We are given that the data set $D = \{(x_i, y_i)\}_{i=1}^m$ satisfies $y_i = \begin{cases} 1 & \text{if } w^T x_i + \theta \geq 0 \\ -1 & \text{if } w^T x_i + \theta < 0 \end{cases}$

Consider the linear program

$$\begin{aligned} \min & \delta \\ \text{subject to} & y_i(w^T x_i + \theta) \geq 1 - \delta \\ & \delta \geq 0 \end{aligned}$$

(b) Solution: Given there is an optimal solution with $\delta = 0$ then $y_i(w^T x_i + \theta) \geq 1$ is satisfied for all D and

$$\begin{aligned} \text{subject to} & y_i(w^T x_i + \theta) \geq 1 \geq 0 \text{ when } y = 1 \\ & y_i(w^T x_i + \theta) \leq -1 < 0 \text{ when } y = -1 \end{aligned}$$

which shows it to be linearly separable.

(c) Solution: If there is a hyperplane that satisfies $y_i(w^T x_i + \theta) \geq 1 - \delta$ for $\forall D$ with $\delta > 0$ then the value of delta determines our linear separability. if $\delta < 1$ then we know its separable by (b), if $\delta > 1$ if we cannot tell if the data set is separable, and if $\delta = 1$ we know if is not separable.

(d) Solution: The trivial optimal solution is $w = 0$, $\theta = 0$, and $\delta = 0$. This is a solution because $y_i(w^T x_i + \theta) \geq -\delta$ is satisfied. The issue with this formulation is that the solution doesn't form a hyperplane.

(e) Solution: The possible optimal solutions are since there are only two points in our D we know the dataset is linearly separable. Plugging into the linear program it becomes apparent we need to satisfy

$$\begin{aligned} 1 \cdot (w_1 + w_2 + w_3 + \dots + w_n + \theta) &\geq 1 \\ -1 \cdot (-w_1 - w_2 - w_3 - \dots - w_n + \theta) &\geq 1 \end{aligned}$$

which means we need $w_1 + w_2 + w_3 + \dots + w_n \geq |1 + \theta|$

so the set of optimal solutions are all δ, w, θ such that $\delta = 0$, and $w_1 + w_2 + w_3 + \dots + w_n \geq |1 + \theta|$.

Problem 5

(a) Solution: For the training data I see that the data has a negative sloped linear shape. I feel a negative sloped line would make a decent linear regression for predicting this data.

For the testing data I see two main groups around (0.2, 1.5) and (0.8, 0.4) I do not feel like a linear regression could do a good job predicting this data.

(b) Solution: I modified the code in regression.py as instructed

(c) Solution: Completed PolynomialRegression.predict as required

(d) Solution:	rate	coefficients	iterations	final value
	10^{-4}	[1.91573, -1.743589]	10000	5.49356
	10^{-3}	[2.44638, -2.81630]	10000	3.91257
	10^{-2}	[2.44640, -2.81635]	1505	3.91257
	0.0407	[2.44640, -2.81635]	381	3.91257

The coefficients all seem to converge to [2.44, -2.81] which is good, means that there is consistency so they are probably converging correctly. It seems that the larger the η the more efficient the convergence. $\eta = .0407$ converges much faster than $\eta = 10^{-2}$ the final values are approximately the same

(e) Solution: The closed form solution's coefficients are $[2.4464, -2.8163]$ which are approximately the same as those found by our gradient decent algorithm. The cost of the closed form solution was 3.9125 as well. Overall these numbers match the outputs of our gradient decent algorithm. The closed form solution algorithm runs extremely quickly compared to gradient descent as we do not need to loop 10000 times.

(f) Solution: With the learning rate of $\frac{1}{1+k}$ the algorithm takes 10000 iterations and therefore doesn't converge. The final coefficients are still close $[2.44634, -2.81623]$ and the final value is 3.91257 but with the minimal margins of difference ive been seeing with my numbers this was far worse than the others.

(g) Solution: I updated `PolynomialRegression.generate_polynomial_features(...)` to create an $m + 1$ dimensional feature vector for each instance.

(h) Solution: RMSE is a better estimator of error over $J(\theta)$ because it gives a much more linear error correlation between our data points and our model.
I implemented `PolynomialRegression.rms_error(...)`.

(i) Solution: In the graph below you can see that the best degree polynomial is 5. This is because it has some of the lowest error on the training data and the test data. You can see evidence of overfitting on polynomials of degree 9 and 10. These two polynomials have very low training error because the complex polynomial is fitting their data points more perfectly, however the error for the test data blows up because of the same complex polynomial is very bad at fitting new points.

